

Part 1: Introduction

- **What is an Operating System (OS)?**
- Types of Computing Systems
- Computer System Architecture (Review)
- Operating System Services

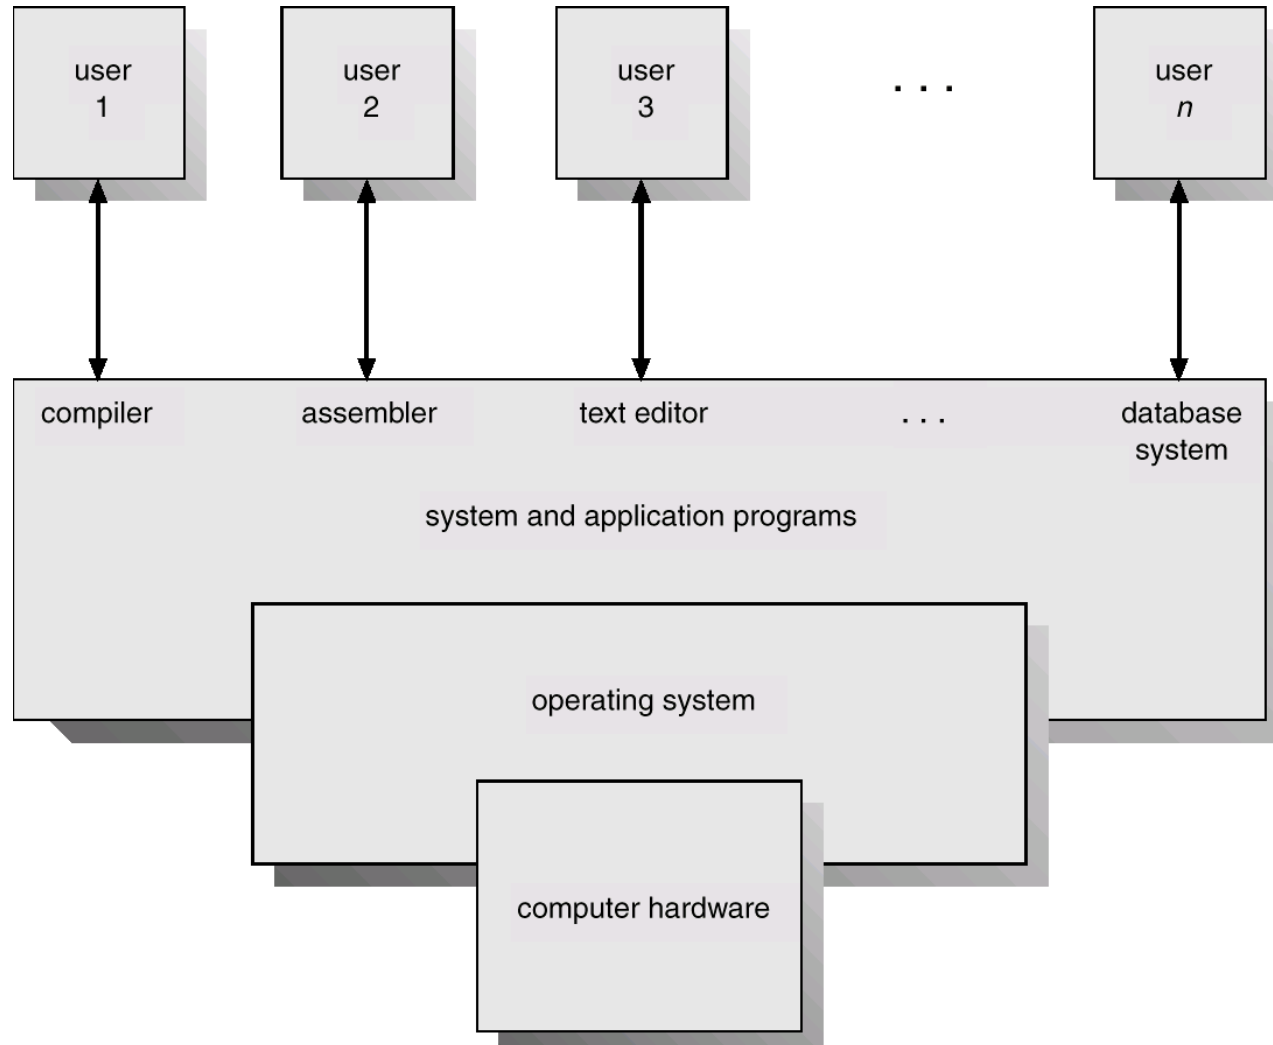
What is an Operating System?

- A **program** that acts as an **intermediary** between a user and the computer hardware
- **Two major goals:** user convenience and efficient hardware utilization
 - Hide hardware complexity
 - Use hardware in an efficient manner
 - Smart resource (e.g., Central Processing Unit (CPU), I/O devices, memory) allocation
- **These two goals can be contradictory**
 - Smart resource allocation may require lot of information about user programs

Computer System Components

1. **Hardware:** Provides basic computing resources (CPU, memory, I/O devices)
2. **Operating System(OS):** Controls and coordinates the use of the hardware among various application programs for various users
3. **Application programs:** Defines the ways in which system resources are used to solve computing problems for users (compilers, database systems, video games, business programs)
4. **Users:** People, machines, other computers

Abstract View of System Components



Operating System Definitions

- **Resource allocator:** Manages and allocates hardware resources
- **Control program:** Controls the execution of user programs and operations of I/O devices
- **Kernel:** The one “core” program that is always ready to accept new commands from the users as well as hardware (all else being application programs)

Part 1: Introduction

- What is an Operating System (OS)?
- **Types of Computing Systems**
- Computer System Architecture (Review)
- Operating System Services

Types of Computing Systems

- Batch systems
- Multiprogrammed and Time-sharing systems
 - Desktop systems
- Embedded and Cyber-physical systems
 - Real-Time systems
 - Handheld systems

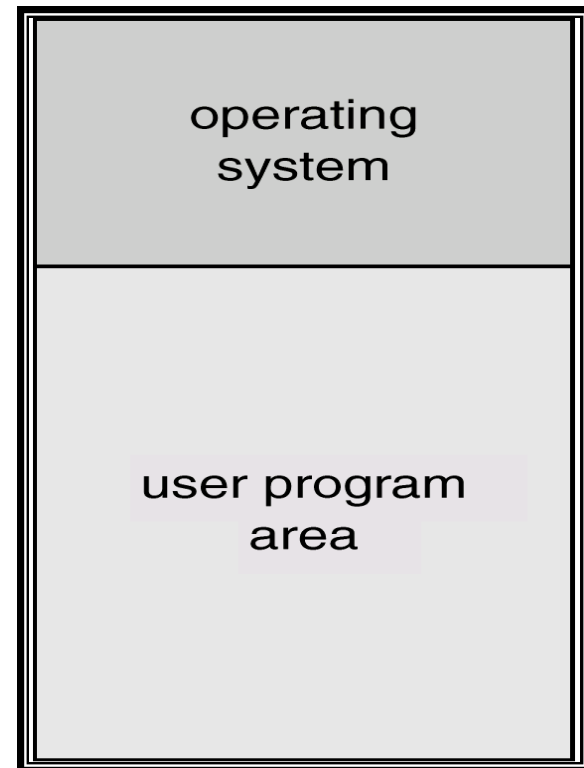
A Short Review on Computers

In this video, we will briefly review the history of computers




Simple Batch Systems

- Reduce setup time by batching similar jobs
- **Automatic job sequencing** – automatically transfers control from one job to another
- **Simple memory layout**
 - Only one user job in memory at any time point
- **Not very efficient**
 - When job waits for I/O, CPU idles



Multiprogrammed (Time-Sharing) Systems

- Several jobs are kept in main memory at the same time, and the CPU is multiplexed among them
- A job is **swapped** in and out of memory to the hard disk
- System is highly interactive; supports multiple online users
- Examples: desktops, servers



OS Features Needed for Multiprogramming

- **Memory management:** To allocate memory to several jobs
- **CPU scheduling:** To choose among several jobs ready to run
- **I/O device scheduling:** Allocation of I/O devices to jobs

Desktop Systems

- Personal computers – computer system dedicated to a single user
- Several I/O devices – keyboard, mouse, display screens, printers, etc.
- User convenience and responsiveness is the main focus
- May run several different types of operating systems (Windows, MacOS, UNIX, Linux)

Embedded and Cyber-physical Systems

- **What are they?** Physical systems whose operations are monitored and controlled by a reliable computing and communication core
- **Resource constrained:** Low power, small memory, low bandwidth, etc.
- **Domain-specific OSes:** Real-time, Handheld, Automotive, Avionics, Industrial Controls, Sensor networks, etc.



Real-Time Systems

- Used as a control device in a dedicated application such as industrial controls, automotive, avionics, medical devices, etc.
- **Well-defined fixed-time constraints**
 - Job must be completed within a deadline
 - Example: Airbag control in cars
- Example real-time OSes




Handheld Systems

- Mobile phones, tablets
- Issues:
 - Limited memory
 - Slow processors
 - Small display screens
- Popular OSes: iOS, Android, Windows Phone



A Dual-Core CPU Design



Multiprocessor Systems

- Systems with more than one CPU, or CPU with multiple cores (also called multi-core systems)
- *Tightly coupled systems*: Communication usually takes place through shared memory
- **Advantages of such parallel systems**
 - Increased system *throughput*
 - Economical *due to* sharing of memory and I/O devices (as compared to multiple single CPU systems)
 - Increased reliability due to redundancy

Part 1: Introduction

- What is an Operating System (OS)?
- Types of Computing Systems
- **Computer System Architecture (Review)**
- Operating System Services

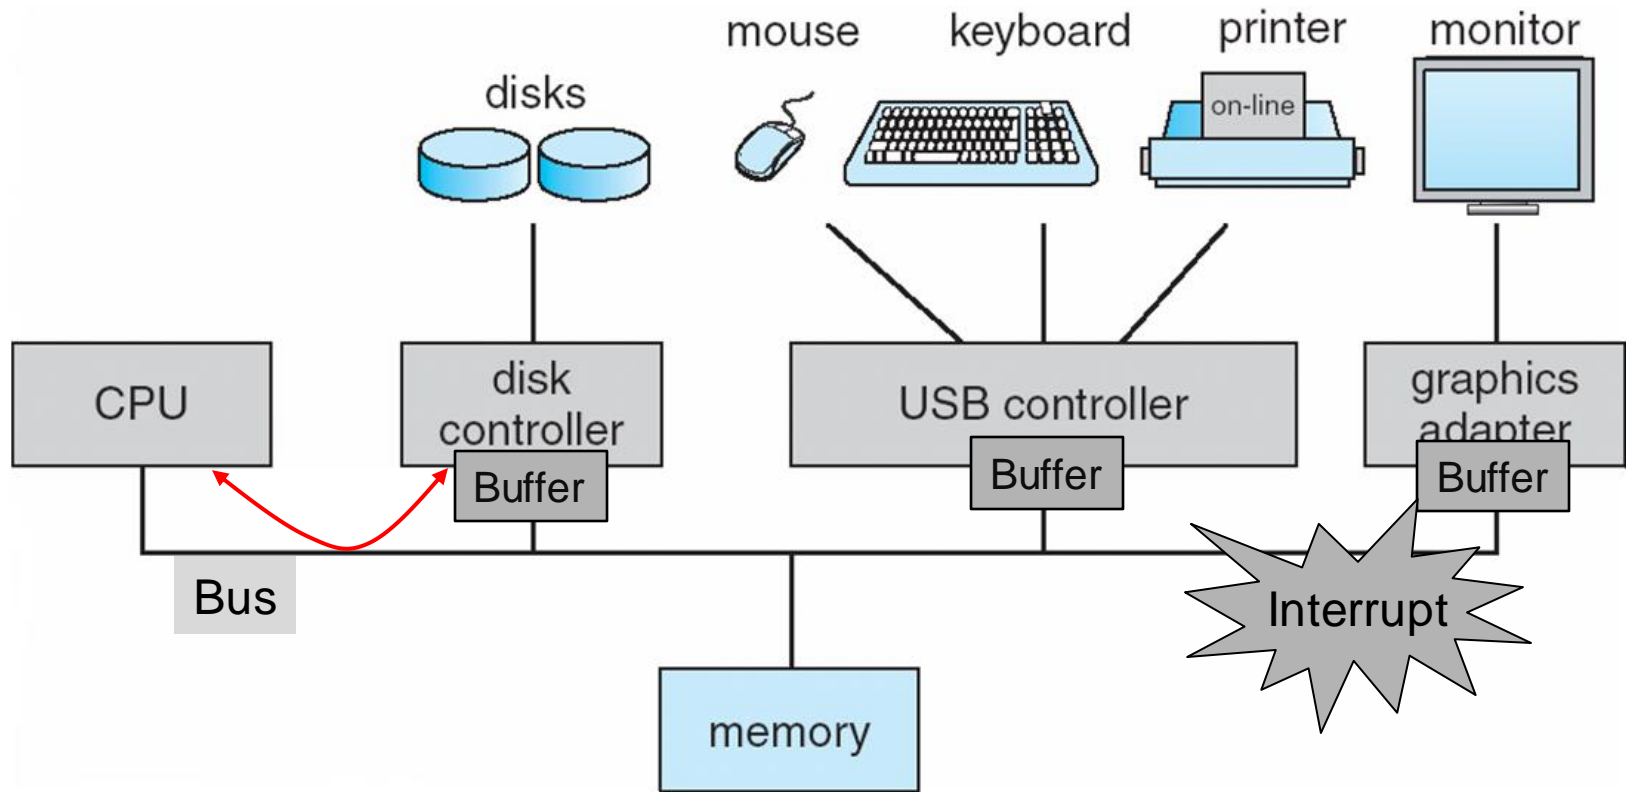
Computer System Architecture

- Computer System Operation
- Storage Hierarchy
- Hardware Protection

Computer-System Operation

- I/O devices and the CPU can **execute concurrently**
- Each **device controller** is in charge of a particular device type
- Each device controller has a local **buffer**
- Device controller moves data between local buffer and memory
- Device controller informs CPU that it has finished its operation by causing an ***interrupt***

Computer-System Operation




Common Functions of Interrupts

- An interrupt transfers control to the *interrupt service routine* generally through the *interrupt vector*, which contains the addresses of all the service routines
- Incoming interrupts are *disabled* while another interrupt is being processed to prevent any *loss of interrupts*
- A *trap* is a CPU generated interrupt caused either by a software error or request
 - Unhandled exceptions in user program
- An operating system is typically *interrupt driven*
 - If the OS is not interrupt driven, it would be required to poll for task/event completions

Interrupt Handling

- The operating system preserves the state of the CPU by storing registers and the program counter
 - Also called a *context switch*
- It then determines which type of interrupt has occurred
 - Separate segments of code determine what action should be taken for each type of interrupt
- Based on the interrupt type, it identifies the appropriate *interrupt service routine* to execute
 - Obtained from the *interrupt vector table*


Interrupt-Driven I/O Timeline



Direct Memory Access (DMA)

- Used for **high-speed I/O devices** that are able to transmit information at close to memory speeds
- OS sets up the memory blocks, counters, etc.
- Device controller transfers data blocks from buffer to main memory **without CPU intervention**
- Only **one interrupt is generated per block**, rather than one interrupt per byte

How a Modern Computer Works




Storage Hierarchy

- **Memory Hierarchy:** CPU registers, CPU caches, main memory, hard disk ...
- Storage system organization based on
 - speed
 - cost
 - volatility
 - size
- **Caching:** Copying information into faster storage system; main memory can be viewed as the last *cache* before secondary storage

This slide is not examinable.

Storage Hierarchy



Size increases
Speed decreases
Cost/price decreases

This slide is not examinable.

Hardware Protection


- Dual-Mode Operation
- I/O Protection
- Memory Protection

Dual-Mode Operation

- Provides hardware protection by differentiating between at least two modes of operations
 1. *User mode*: Execution of user processes
 2. *Monitor mode (supervisor mode or system mode or kernel mode)*: Execution of operating system processes

Dual-Mode Operation (Cont.)

- *Mode bit* added to computer hardware to indicate the current mode: monitor (0) or user (1)
- When an interrupt or trap occurs hardware switches to monitor mode



- *Privileged instructions can be used only in monitor mode*

Kernel mode vs. root/admin.

- Are they the same?
 - No, they are not the same
- Kernel or user mode is a hardware operation mode
- Root/Administrator is a user account in an OS
 - Jobs still execute in user mode, even when executed by a root/admin. user
 - This user may execute code in kernel mode indirectly → e.g., by loading a kernel module

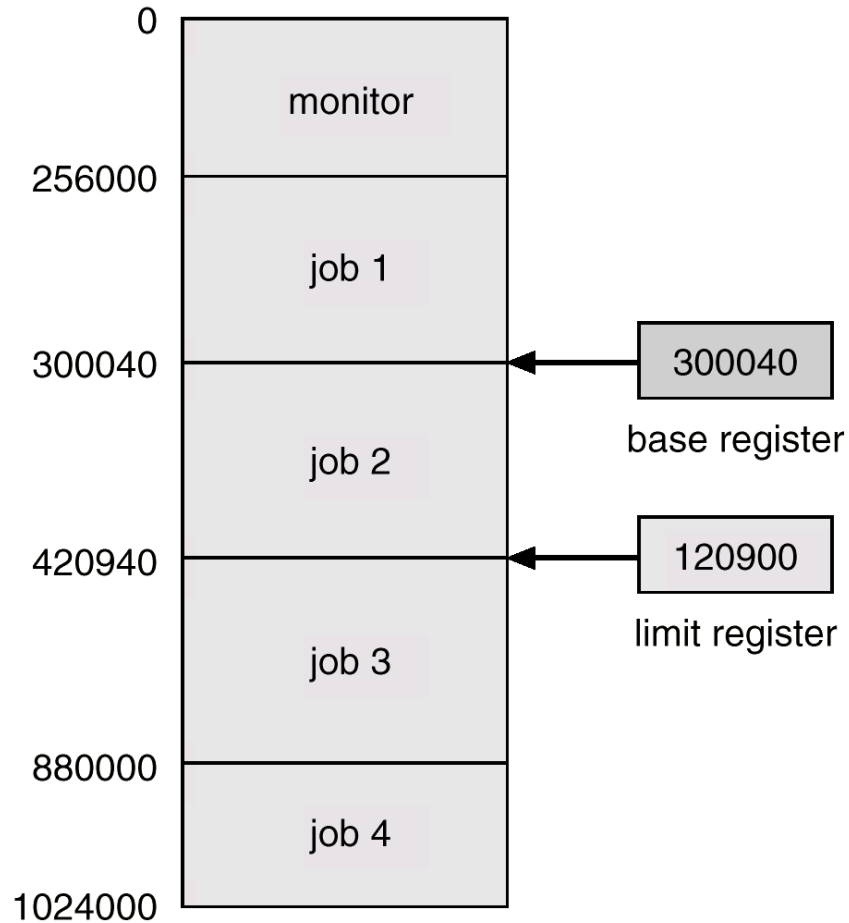
I/O Protection

- A user program may issue illegal I/O operation; hence I/O must be protected
 - Case 1: read a file that does not exist
 - Case 2: unauthorized access to a device
- All I/O instructions are privileged instructions
- All I/O operations must go through the OS to ensure its correctness and legality
 - CPU generates a trap for I/O operations that try to bypass the OS

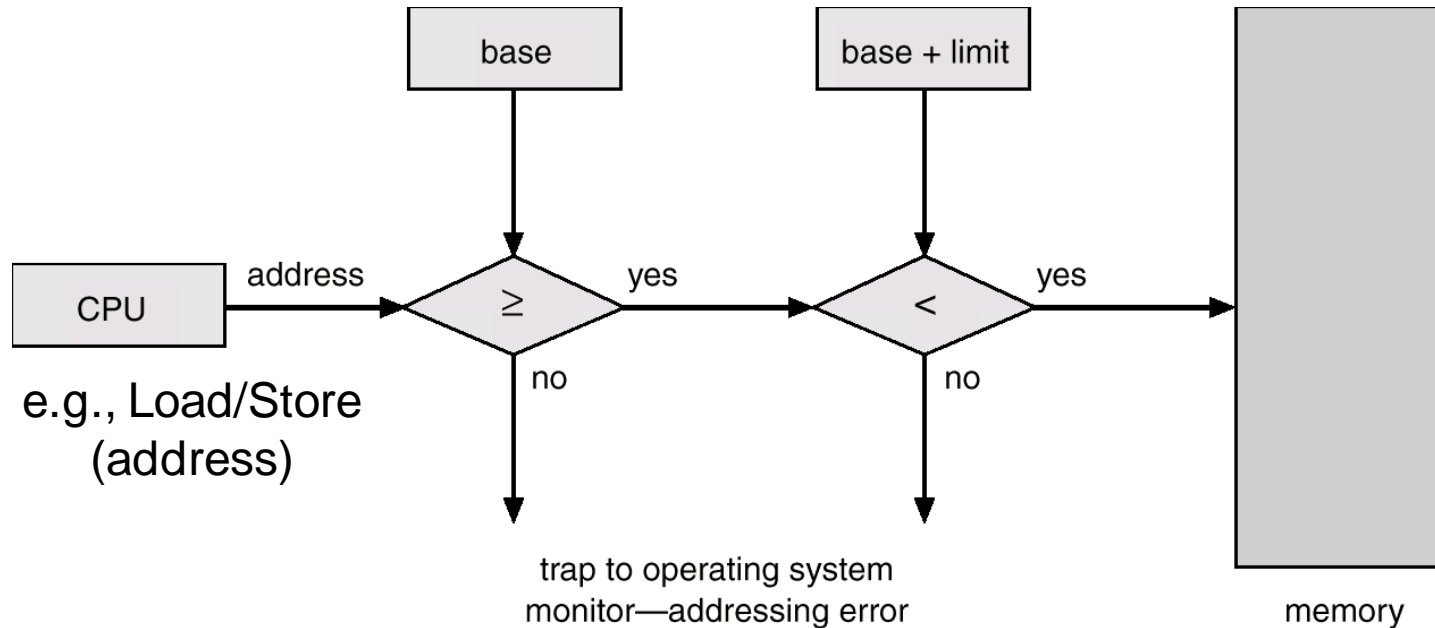
Memory Protection

- Must provide memory protection, at least for the interrupt vector and the interrupt service routines
- Two CPU registers determine the range of legal addresses a program may access:
 - **Base register:** Holds the first legal physical memory address
 - **Limit register:** Contains the size of legal range
- Memory outside the defined range is protected and cannot be accessed

Memory Protection (Cont.)



Memory Protection (Cont.)

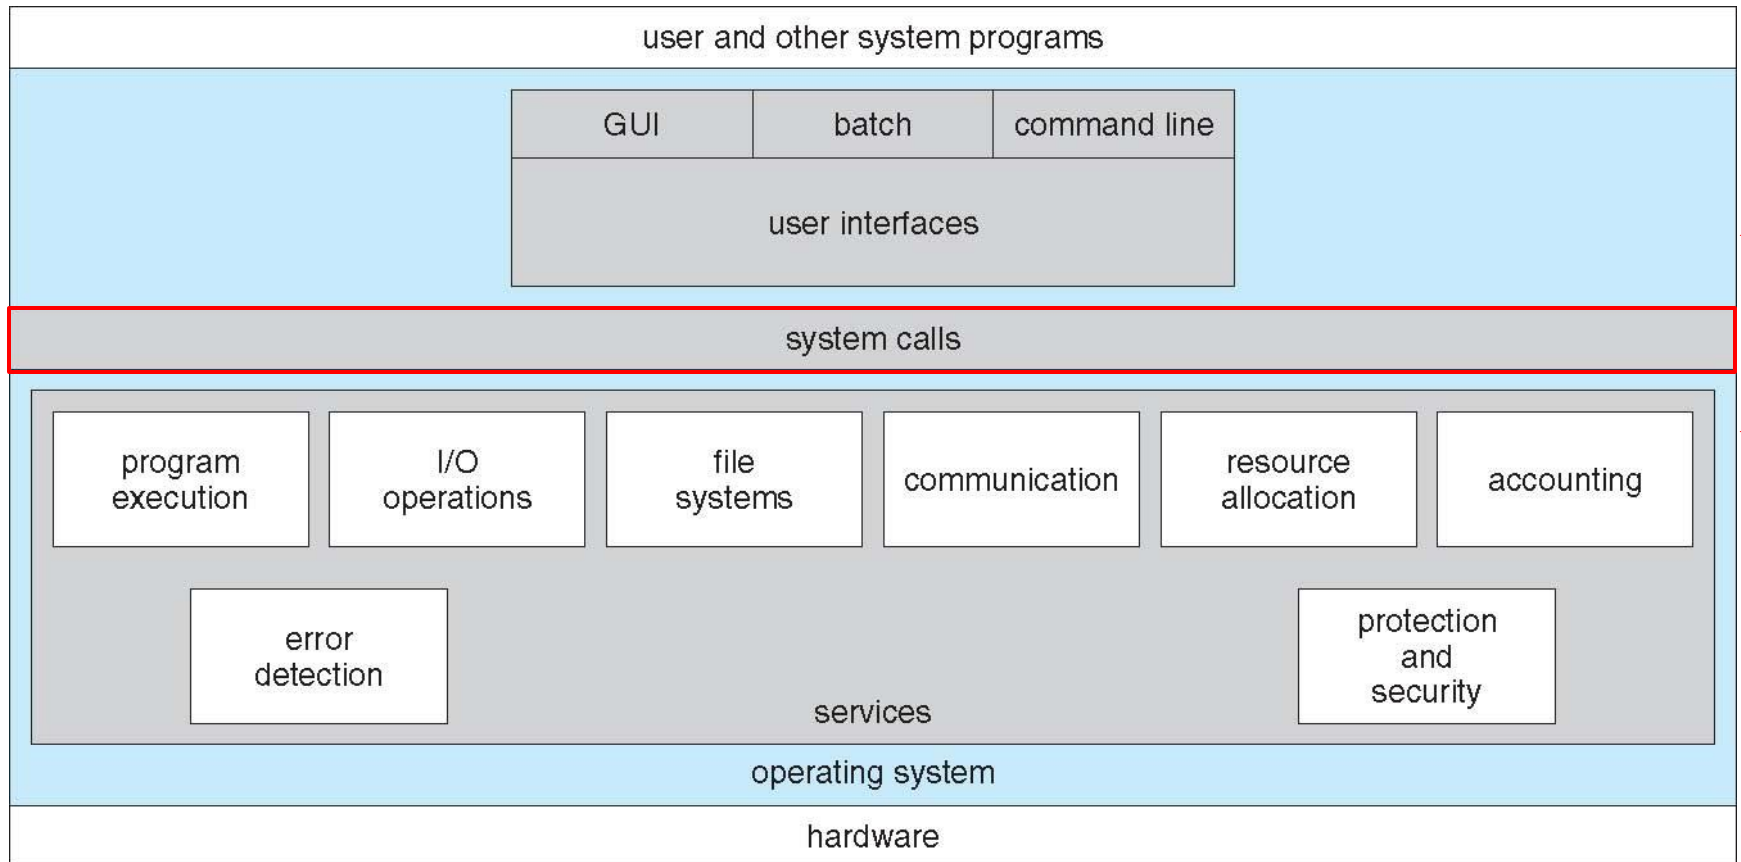


- The load instructions for the *base* and *limit* registers are privileged instructions (only in monitor mode)
- CPU issues a trap to the OS if above checks fail

Part 1: Introduction

- What is an Operating System (OS)?
- Types of Computing Systems
- Computer System Architecture (Review)
- **Operating System Services**

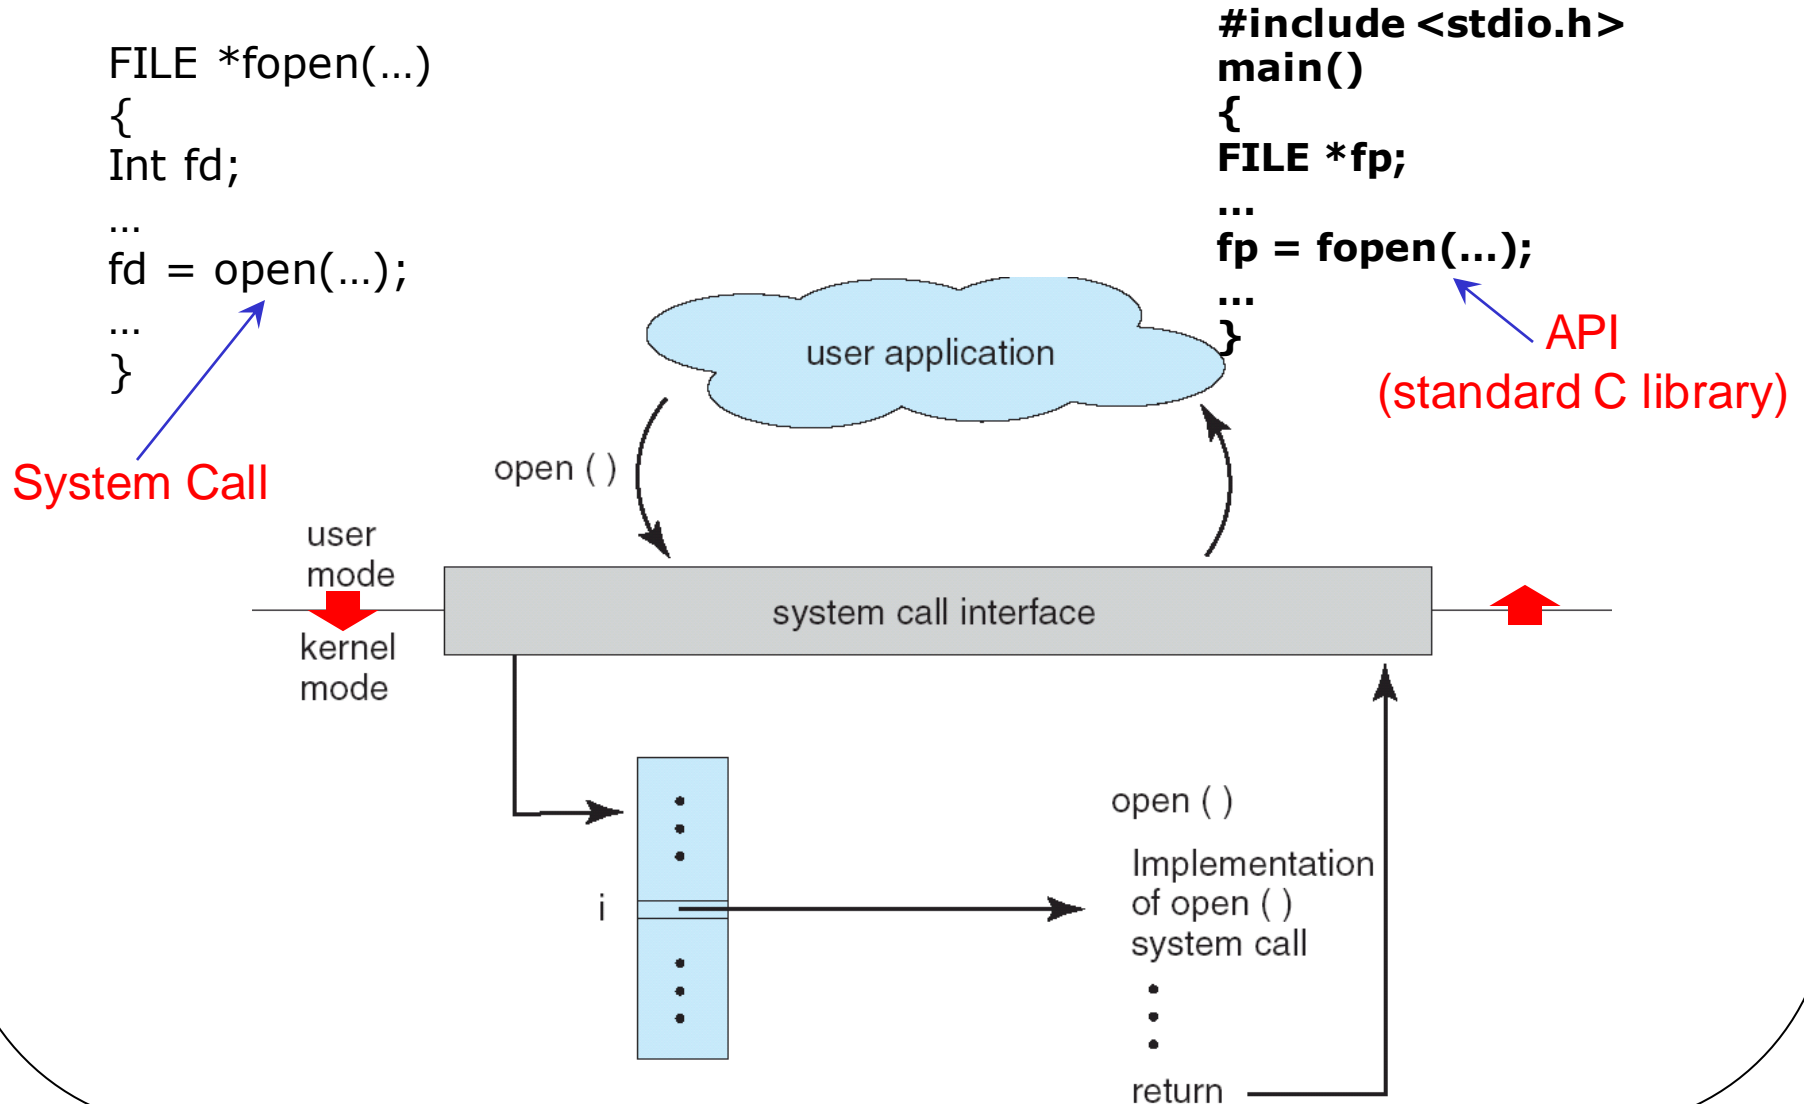
A View of Operating System Services



System Calls

- System calls provide the interface between a user program and the operating system
 - Generally available as assembly-language instructions
 - Possible to replace assembly language for systems programming to allow system calls to be made directly (e.g., in C/C++)
- The execution of a system call requires the switch from the user to the kernel mode

API – System Call – OS Relationship



Advanced Readings

- Mobile operating systems
 - Apple iOS, <http://en.wikipedia.org/wiki/IOS>
 - Google Android,
[http://en.wikipedia.org/wiki/Android_\(operating_system\)](http://en.wikipedia.org/wiki/Android_(operating_system))
 - Microsoft Windows Phone,
http://en.wikipedia.org/wiki/Windows_Phone
- “History of Operating Systems”, by Ayman Moumina (pdf in NTULearn)

Future Operating Systems

- Operating systems are **not** limited to Windows/Linux/macOS etc.
- Operating system development never stops
 - Changes in hardware
 - Changes in user requirements
- Examples:
 - Microkernel (also known as μ -kernel), <https://en.wikipedia.org/wiki/Microkernel>
 - Formally verified OS kernel, <https://sel4.systems/>