

Part 6: Real-Time OS & Virtualization

- **What is a Real-Time OS (RTOS)?**
- Real-Time Process Specification
- Real-Time CPU Scheduling
- Virtualization


Cyber-Physical Systems

- Physical/Engineered systems whose operations are monitored, coordinated and controlled by a reliable computing and communication core
 - Automotive Systems (Autonomous driving, Parking assist, Airbag controls)
 - Avionics (Flight navigation & control)
 - Manufacturing Systems (Robotics, Process controls)
 - Medical Systems (Robotic surgery, devices)
 - ...


Relevance of Real-Time

- Common to many application examples we saw in the previous slide:
 - **Collect data** from various sensing devices
 - **Execute control law(s)** to determine response
 - **Send actuator commands** in a **reasonable amount of time**

Pacemaker timing diagram



Collision avoidance and braking



Collision Warning with Auto Brake

What is a Reasonable Time?

- What is the functionality?
 - Collision avoidance in automotive (milliseconds)
 - Pacemaker (up to a second)
 - Robotic surgery (varies greatly depending on the target)
- What are the environment constraints?
 - Available computing and communication resources
 - Timing characteristics of sensors/actuators/operations
- Failure-mitigation strategies?
 - Time to detect and recover from failures
 - Example: execution replication for redundancy

Common Misconception

- Real-Time \neq Fast
- Real-Time = Predictable even in the worst-case



“Man drowned in a river with average depth 20 cms”

Real-Time CPS / Real-Time OS


- **Definition:** System whose correctness depends not only on the **logical/functional aspects**, but also on the **temporal aspects**
 - Application has deadlines that must be met
 - **A real-time OS (RTOS) provides OS services to such systems (e.g., FreeRTOS, MicriumOS, ...)**
- **Key performance measure for RTOS**
 - Timeliness/Predictability on timing constraints (**deadlines**)
 - Significance of **worst-case** over average-case
 - Deadlines are a function of application requirements

Part 5: Real-Time OS & Virtualization

- What is a Real-Time OS (RTOS)?
- **Real-Time Process Specification**
- Real-Time CPU Scheduling
- Virtualization

RTOS (Real-Time) Process

- **Definition:** A real-time process is specified as $\langle R, C, D \rangle$, where R is process release time, C is execution requirement and D is relative deadline
 - Requires C time units of CPU in the interval $[R, R+D)$
 - **How does one determine these parameters?**
- Example: $P_1 \langle 0, 5, 10 \rangle$, $P_2 \langle 5, 10, 35 \rangle$, $P_3 \langle 20, 10, 30 \rangle$




In this lecture, we assume processes only have a single CPU burst; C is the duration of this burst

Recurrent Real-Time Process

- **Nature of real-time processes**
 - Collect data from sensing devices, execute control laws to determine responses, and send actuator commands in reasonable time
 - **Repeat the above steps regularly**
 - * Examples: airbag control, flight control, collision avoidance, pacemaker, etc.
- **A recurrent real-time process**
 - **Executes some function repeatedly over time**
 - Each instance of execution is a real-time process $\langle R, C, D \rangle$


Periodic Real-Time Process

- **Definition:** A process that **repeats periodically**
 - Processes generated by a **time-triggered phenomena** (sensor sending data periodically)
 - Example: Perception function for collision detection
- A periodic process is specified as $\langle T, C, D \rangle$, where T is process period, C & D are as defined earlier
 - Real-time processes are released at $R=0, T, 2T, \dots$
 - Example: Periodic process $P\langle 10, 5, 7 \rangle$



Sporadic Real-Time Process

- **Definition:** A process that **repeats sporadically with a minimum gap between releases**
 - Processes generated by an **event-triggered phenomena**
 - Example: Anti-lock braking function in automotive
- A sporadic process is specified as $\langle T, C, D \rangle$, where T is minimum release-separation time
 - Real-time processes are released with a min. gap of T
 - Example: Sporadic process $P\langle 10, 5, 7 \rangle$




Part 5: Real-Time OS & Virtualization

- What is a Real-Time OS (RTOS)?
- Real-Time Process Specification
- **Real-Time CPU Scheduling (short-term scheduler)**
 - **Fixed-priority scheduling**
 - **Dynamic-priority scheduling**
- Virtualization


Why Classic Algorithms Fail?

- Consider real-time processes (non-recurring): $P_1<0,20,30>$, $P_2<0,5,10>$, $P_3<5,10,35>$, $P_4<20,10,30>$, $P_5<25,5,10>$


FCFS schedule



SJF schedule



RR(q=5) schedule

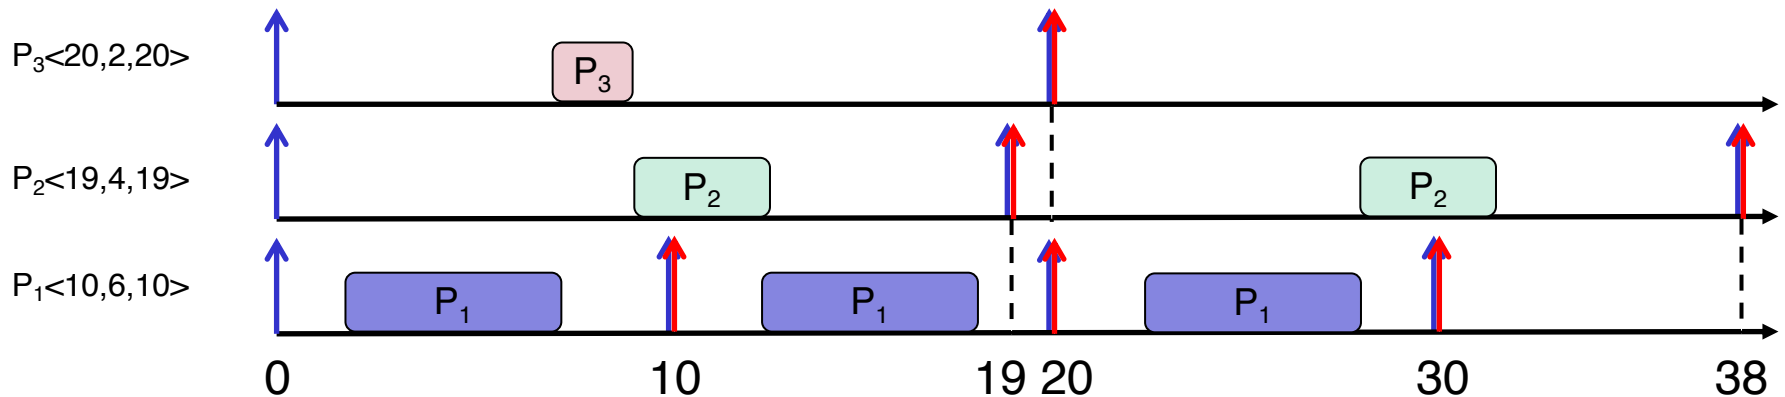


They don't prioritize deadlines and hence perform poorly

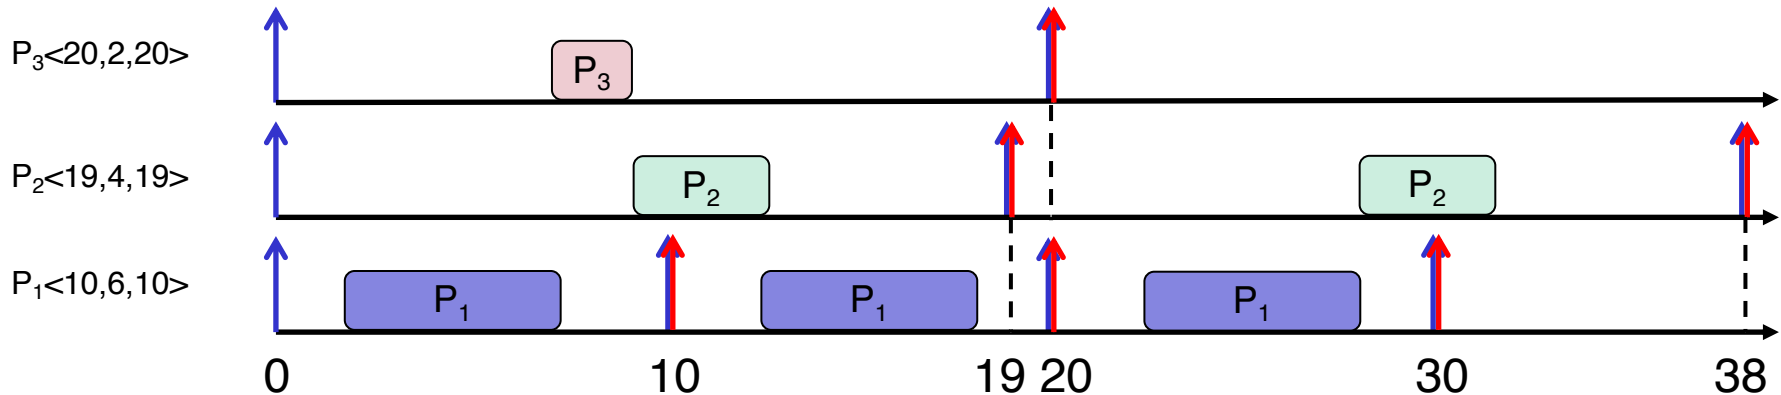
Real-Time CPU Scheduling Problem

Given a set of periodic/sporadic real-time processes, find a uni-processor CPU scheduling algorithm that can meet process deadlines

- We will use a running example (periodic real-time process set): $P_1<10,6,10>$, $P_2<19,4,19>$, $P_3<20,2,20>$



Fixed-Priority CPU Scheduling




Priorities are fixed for each recurrent process

- **Priorities are fixed across instances of recurrent processes**
- Suppose instance of P_1 ($R=0$) has higher priority than instance of P_2 ($R=0$). Then,
 - * P_1 ($R=10$) has higher priority than P_2 ($R=0$)
 - * P_1 ($R=10$) has higher priority than P_2 ($R=19$)
 - * P_1 ($R=20$) has higher priority than P_2 ($R=19$) ...

Rate Monotonic (RM) Scheduler

- Assign priorities based on process periods / minimum release-separation time (T)
 - Shorter T implies higher priority
 - Ties are broken arbitrarily



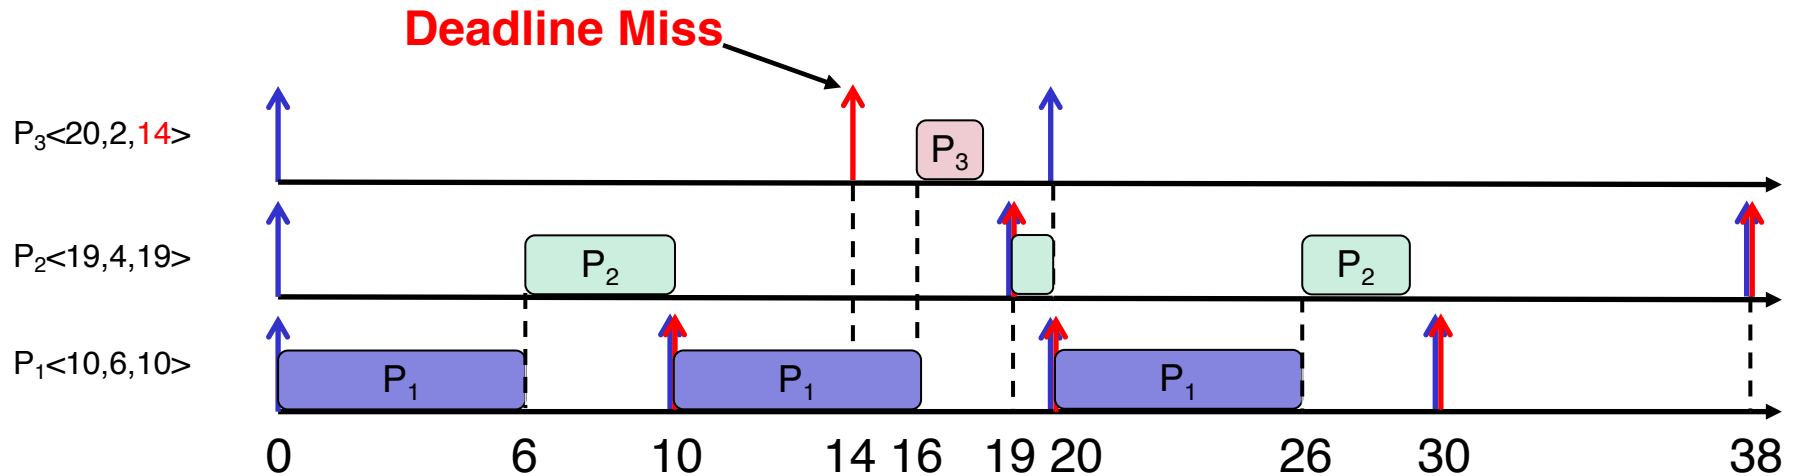
RM is a very popular short-term CPU scheduler in the real-time CPS industry. Why?

RM and Process Deadlines

- RM is good, but still does not always prioritize urgent processes


– Suppose we modify the process set as follows:

$P_1\langle 10, 6, 10 \rangle$, $P_2\langle 19, 4, 19 \rangle$, $P_3\langle 20, 2, 14 \rangle$



Deadline Monotonic (DM) Scheduler

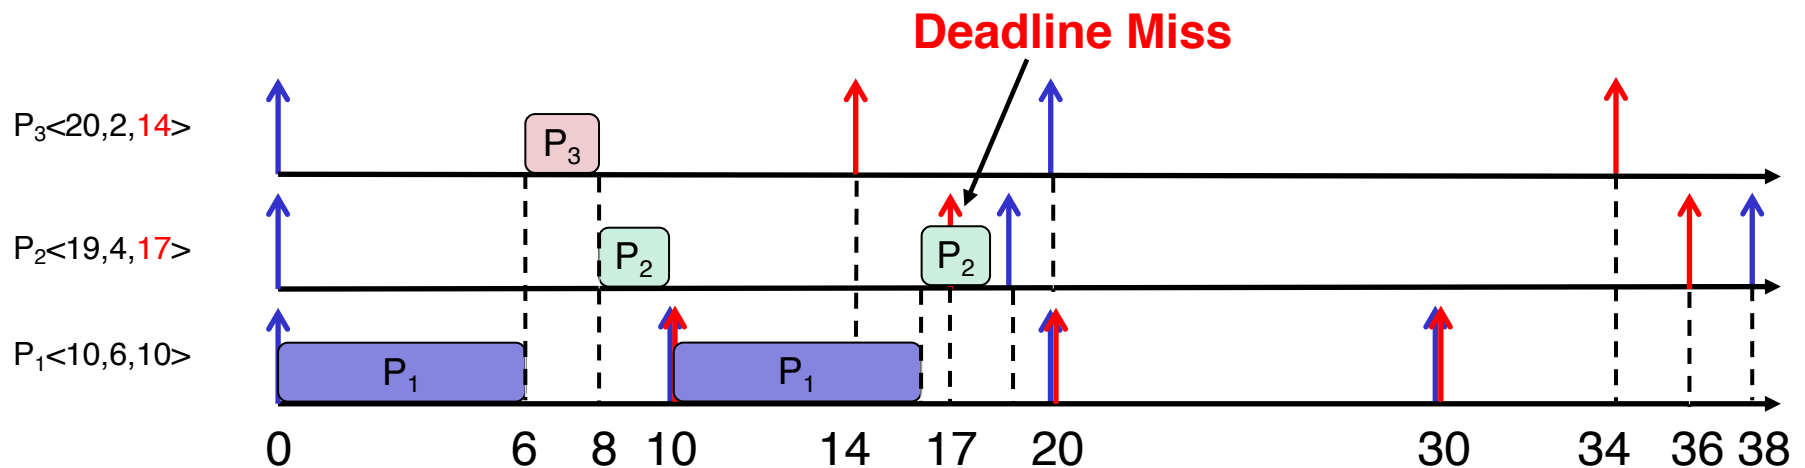
- Assign priorities based on process deadlines (**D**)
 - Shorter **D** implies higher priority
 - Ties are broken arbitrarily



Both RM and DM are fixed-priority schedulers


DM and Process Deadlines

- DM is better than RM, but it cannot change priorities across process instances
 - Suppose we further modify the process set as follows:
 $P_1\langle 10, 6, 10 \rangle$, $P_2\langle 19, 4, 17 \rangle$, $P_3\langle 20, 2, 14 \rangle$



Earliest Deadline First (EDF) Scheduler

- Dynamic-priority scheduler that assigns priorities based on process **instance deadlines**
 - Instances with shorter deadline are given higher priority
 - * **NOT the same as parameter D**
 - Ties are broken arbitrarily



EDF is a dynamic-priority scheduler, hence more powerful than RM and DM

RM/DM versus EDF


RM/DM

- **Simpler implementation**
(separate queue for each recurrent process)
- **Predictability for high priority processes, even under overload**
(next slide)

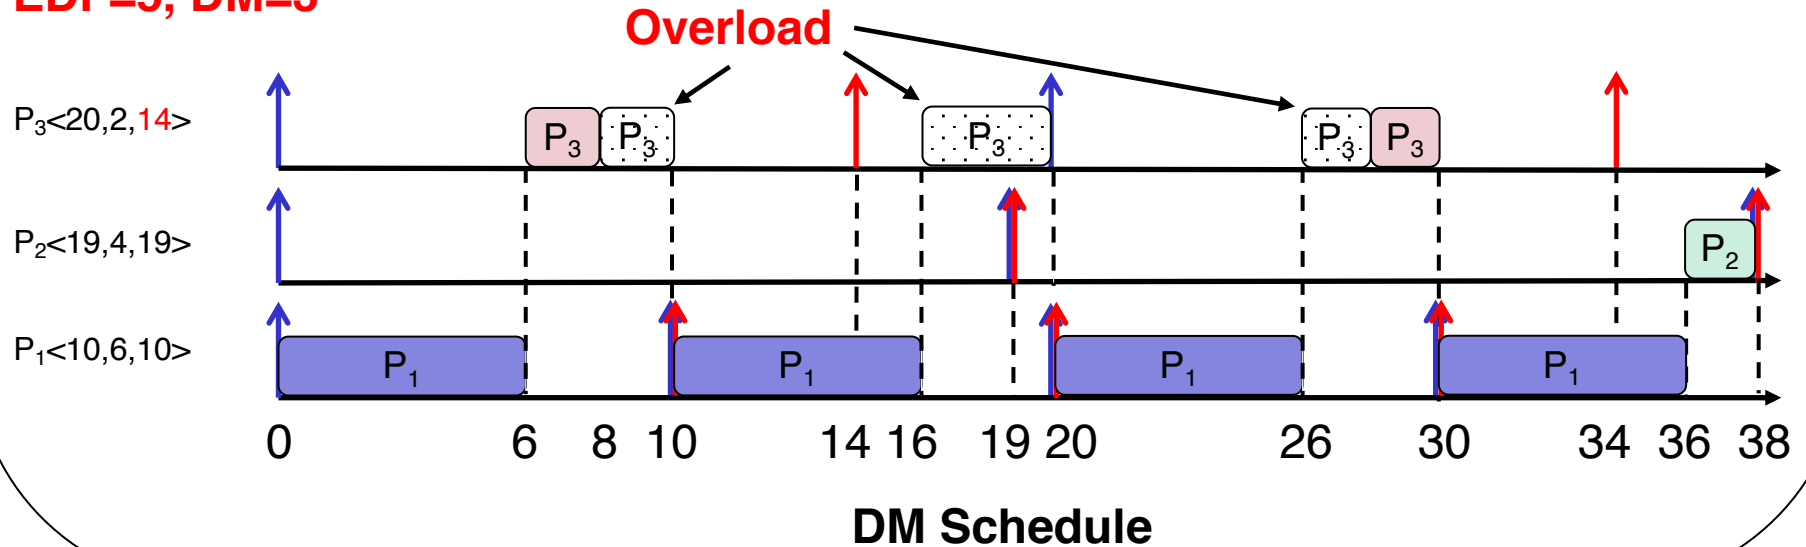
EDF

- **Harder implementation**
(online sorting of queue based on instance deadlines)
- **Misbehaviour during overload**
(next slide)

DM versus EDF during Overload



Deadline Misses
EDF=5, DM=3



Part 5: Real-Time OS & Virtualization

- What is a Real-Time OS (RTOS)?
- Real-Time Process Specification
- Real-Time CPU Scheduling
- **Virtualization**

What is Virtualization?

- **Definition:** It is a technique that uses software, called **Hypervisor (virtual machine manager or VMM)**, to create abstraction of hardware
 - Hardware is divided into multiple virtual computers, called **Virtual Machines (VMs)**
 - Each VM runs its own OS, called **Guest OS**, and behaves like an independent computer
 - * Application processes can run on the guest OS as if it is an independent computer
 - Each VM is using only a portion of the actual hardware

Is a general OS also using virtualization?

Functions of a Hypervisor

- **Creation and management of VMs**
 - Allocating hardware resources for VMs
 - Executing instructions on the hardware on behalf of VMs (VMs may be using different guest OSes)
- **Communication between VMs**
 - Mechanisms for VMs hosted on the same hardware to be able to communicate securely with each other
- **VM Migrations**
 - Migrating VMs from one hypervisor/hardware to another, almost instantaneously (at runtime)
 - Gives a lot of flexibility and portability

Why do we need Virtualization?

- **Allows for more efficient utilization of hardware**
 - Cost-effective hardware deployment & sharing
 - Low latency and agile execution-environment deployments (VM creation and flexibility)
 - Failure mitigation (VM independence and migrations)
- **Key technology that is driving cloud computing**
 - Cloud providers can dynamically and cost-effectively scale hardware allocations based on user requirements
 - Enables the concept of platform as a service – rent hardware & services as and when needed

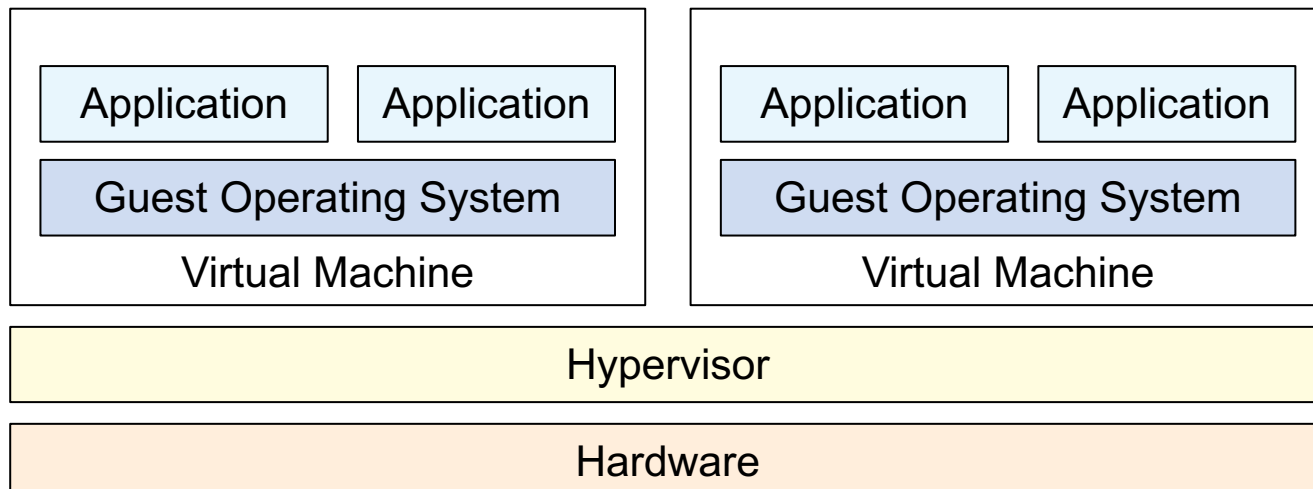
Virtualization Challenges

- Requires management layer: Hypervisor
 - Hypervisor is usually two orders of magnitude smaller than general purpose OS
 - Requires more disk space and RAM
 - Must ensure that instructions can be executed on H/W
- Real-Time CPS require different solutions than server virtualization
 - Real-time (worst-case timing predictability)
 - Minimal memory footprint & minimal overhead (highly resource constrained)

Type-1 Virtualization

Hypervisor interacts directly with hardware

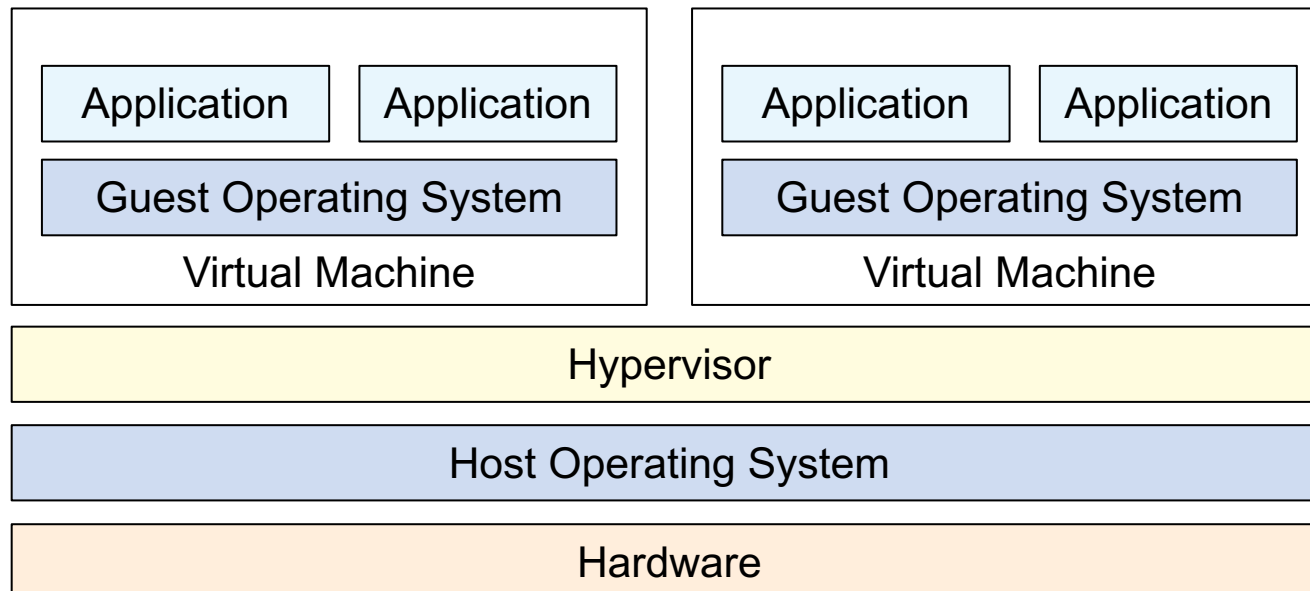
- Also called bare-metal hypervisor
- Highly secure and has low latency
- Popular in industry (KVM, Microsoft hyper-v, VMware esxi, Xen)



Type-2 Virtualization

Hypervisor interacts with a host OS, which in turn interacts with hardware

- Also called hosted hypervisor
- Higher latency, less popular, mostly for end-user virtualization (Oracle virtualbox, VMWare workstation)



Virtualization Levels

1. Full virtualization

- Complete abstraction of the actual hardware
- Guest OS runs unmodified on the hypervisor
 - * It is unaware of hypervisor's existence
- Examples: VMware esxi, Microsoft virtual server

2. Para-virtualization

- Unique software interface (API) between guest OS - VM
- Guest OS needs to be modified to adapt to the new API
- **Advantage:** VM does not virtualize hard-to-implement parts of the H/W instruction set, hence more efficient
- Examples: Xen, VMware, XtratuM