

# Operating System Security

Asst Prof. Tianwei Zhang  
12 November 2020

# Self-Introduction

## □ Experience

- 2019–present: Assistant Professor @ SCSE, NTU
- 2017–2019: Software Engineer @ Amazon Web Services
  - **Cloud Computing** and **Robotics** platforms
- 2011–2017: PhD @ EE, Princeton University
  - Thesis topic: **cloud computing security**

## □ Research interest

- Design and develop **secure** computer systems
- Principle: enhance system **security** without compromising **performance**, **cost** and **usability**.

# Questions Before Our Course

- ❑ If you download and run a video player app on your computer, can it delete your photos?
- ❑ If you install a maps app on your smartphone, can it steal your privacy information?
- ❑ If you click the attachment of an unknown email, can it destroy your data on the disk?

Yes, if your OS does not have appropriate security protection, e.g., *app isolation*, *permission management*

# Outline

## ❑ System security basis

- Definition
- Security properties

## ❑ Security Protection Stages employed by OS

- Authentication
- Access Control
- Audit

## ❑ Security Threats in OS

## ❑ Hardware Protection

- Integrity verification
- Trusted Execution Environment

# Outline

## ❑ System security basis

- Definition
- Security properties

## ❑ Security Protection Stages employed by OS

- Authentication
- Access Control
- Audit

## ❑ Security Threats in OS

## ❑ Hardware Protection

- Integrity verification
- Trusted Execution Environment

# OS Becomes more Complex

## □ From single-user to multi-user

- DOS is truly single user
- MacOS, Linux, NT-based Windows are multi-user, but typically only 1 user in PCs.
- Cloud computing allows multiple users all over the world to run on the same OS, and they do not know each other.
- Tradeoff: efficiency versus security

## □ From trusted apps to untrusted apps

- Simple real-time systems: only run one specific app
- Runs verified apps from trusted parties
- Modern PCs and smartphones: run apps from third-party developers
- Tradeoff: functionality versus security

# Complex OS brings More Challenges

## □ Protecting a single computer with one user is easy

- Prevent everybody else from having access
- Encrypt all data with a key only one person knows

## □ Sharing resources safely is hard

- Preventing some people from reading private data (e.g. grades)
- Prevent some people from using too many resources (e.g. disk space, CPU core)
- Prevent some people from interfering with other programs (e.g. inserting keystrokes / modifying displays)

# OS Responsibility

## Functionalities

- ❑ Support multiple users concurrently
- ❑ Manage multiple apps concurrently
- ❑ Connect to the network
- ❑ Sharing data with different domains


## Security goals

- ❑ Protect users from each other
- ❑ Protect apps from each other
- ❑ Protect the system from the network
- ❑ Secure the data sharing



# What's being protected? Resources

- System is **secure** if resources used and accessed as intended under all circumstances



# Security Properties

## ❑ Confidentiality (C)

- Prevent unauthorized **disclosure** of information
- Sensitive information should not be leaked to unauthorized parties

## ❑ Integrity (I)

- Prevent unauthorized **modification** of information
- Critical system state and code cannot be altered by malicious parties

## ❑ Availability (A)

- Prevent unauthorized **withholding** of information or resources
- The resources should be always available for authorized users

## ❑ Other properties

- Accountability: actions of an entity can be traced and identified
- Non-repudiation: unforgeable evidence that specific actions occur

# Outline

## ❑ System security basis

- Definition
- Security properties

## ❑ Security Protection Stages employed by OS

- Authentication
- Access Control
- Audit

## ❑ Security Threats in OS

## ❑ Hardware Protection

- Integrity verification
- Trusted Execution Environment

# Security Protection from OS

❑ OS is responsible for protecting the apps running on it

- OS controls what users/processes can do



# Authentication

## ❑ How does a computer know if I am a correct user?

- **What you know?** password, PIN, public/private keys...
- **What you have?** smartcard, hardware tokens...
- **Who you are?** biometrics, face recognition, voice recognition...

## ❑ How does the system conduct authentication?

- Compare the input credential with the stored one
  - Password file: **/etc/passwd** for UNIX
- Allow entry when the credential matches
  - Assign the user an identifier: 32bit for UNIX

# Hash Function

## □ A one-way function $f$

- Takes an input  $x$  of arbitrary length, and produces an output  $f(x)$  of fixed length.

## □ Pre-image resistant

- Given an input  $x$  it is easy to compute  $f(x)$ , but given an output  $y$  it is hard to find  $x$  so that  $y = f(x)$


## □ Collision resistant

- It is computationally infeasible to find a pair  $(y_1, y_2)$ , such that  $y_1 \neq y_2$  and  $f(y_1) = f(y_2)$

# Password Storage

## □ Hashed passwords

- Passwords are hashed and stored in a password table
- When a user inputs a password, its hash value is computed and checked against the password table.



# Password Security

## □ Why hashed passwords?

- Insider attack: even the attacker can access the password table, he is not able to recover the password from the hash values.

## □ Dictionary attack


- Hashed passwords, especially for human-generated passwords, are still vulnerable to dictionary attack.
- This exploits weakness in human-chosen passwords, which tend to derive from words in natural languages.
  - Guess some commonly used passwords
  - Compute their hash values
  - Look for the same hash values in the password table



# Preventing Dictionary Attack

## □ Password salting

- A salt is added to a password before applying the hash function
- A salt is a random string
- Each password has its own unique salt. So even the same password will have different hash values
- The salt value is stored along with the hash of the password + salt
- The attacker needs more hash computation to recover passwords



# Password Complexity

## ❑ Set up higher requirements for the password

- Larger space (lower case, upper case, numbers, special symbols...)
- Length
- No consecutive repeated characters; Not in a dictionary

## ❑ Pros:

- Increase the difficulty of password guessing attack

## ❑ Cons:

- Hard to remember, and easy to type wrong
- People may try to choose passwords that are easy to remember: attacker can guess them easily as well
- People may reuse the old passwords
- People write down the passwords.




# Access Control

## □ Security policy

- Specifies (subject, verb, object) triples
- Subject = (user, program) pair
- Verb = action
- Object = resources

## □ Monitor

- Checks whether the action should be allowed



# Access Control Policy

## ❑ Who sets policy?

- Users, with some system restrictions

## ❑ How is access control list stored?

- Sparse matrix (default deny), store as list

## ❑ How is policy enforced?

- OS exposes API to apps, with privileged operations
- Checks ACL when API functions are called

# Authorization

## □ Access Control Matrix

- Each column represents an object
- Each row represents a subject
- The entry shows the allowed verbs.

	/etc	/homes	/usr
Alice	Read	Read	Read Write
Bob	Read Write	Read Write	Read Write
Carl	None	None	Read

# Update Access Matrix

## □ Access Control Changes

- Grant capabilities: the owner of the object can grant rights to other users.
- Revoke capabilities: subjects can revoke the rights from others

## □ Six Commands to Alter the Access Matrix

- **create subject  $s$** : creates a new subject  $s$ .
- **create object  $o$** : creates a new object  $o$ .
- **enter  $r$  into  $M_{s,o}$** : adds right  $r$  to cell  $M_{s,o}$ .
- **delete  $r$  from  $M_{s,o}$** : deletes right  $r$  from cell  $M_{s,o}$ .
- **destroy subject  $s$** : deletes subject  $s$ . The column and row for  $s$  in  $M$  are also deleted.
- **destroy object  $o$** : deletes object  $o$ . The column for  $o$  in  $M$  is also deleted.

# More Representations

## □ Access Control List (ACLs)

- For one object, which subject has accesses to it? (check the column in the Access Matrix)

## □ Capability:

- For one subject, which objects it has capability to access? (check rows in the Access Matrix)

## □ Most systems use both

- ACLs for opening an object (e.g. `fopen()`)
- Capabilities for performing operations (e.g. `read()`)

# Data Sharing

## ❑ Problem: multiple users want to access the same file or data

- Give each user the corresponding permissions.
- When a new user joins, the permissions have to be granted again.
- When permissions are changed, need to alter each user.

## ❑ Solution: group

- Set permissions for the group instead of the user
- A user joining the group will have the corresponding permissions.
- A user quitting the group will loss the corresponding permissions.
- Easier to manage and update.



# Audit Logs

## □ Audit trail

- Recording all protection-orientated activities, important to understanding what happened, why, and catching things that shouldn't

<b>/usr/adm/lastlog</b>	Records the last time a user has logged in; displayed with <b>finger</b>
<b>/var/adm/utmp</b>	Records accounting information used by the <b>who</b> command.
<b>/var/adm/wtmp</b>	Records every time a user logs in or logs out; displayed with the <b>last</b> command.
<b>/var/adm/acct</b>	Records all executed commands; displayed with <b>lastcomm</b>
<b>/var/log/</b>	In modern Linux systems, log files are located in there

# Outline

## ❑ System security basis

- Definition
- Security properties

## ❑ Security Protection Stages employed by OS

- Authentication
- Access Control
- Audit

## ❑ Security Threats in OS


## ❑ Hardware Protection

- Integrity verification
- Trusted Execution Environment

# Privileged Rings

## □ Operating modes

- Kernel mode has the highest privileges, running the critical functions and services
- Entities with the higher privilege levels cannot call the functions and access the objects in the lower privilege levels directly.
  - System call, interrupt, etc.
- Status flag allows system to work in different modes (context switching)



# Malware

## ❑ Software code that maliciously subvert the computer system

- **Virus**: malicious program that causes copies of itself to be created when triggered by the user
- **Worms**: malicious program that causes copies of itself to be created without any user intervention
- **Trojan horses**: appears to do something useful, but masks some hidden malicious activities
- **Rootkit**: hides other malware from detection and maintains root-level access to the computer.
- **Backdoor**: allow a remote party to gain access to the computer
- **Bot**: inserted into a computer and lies dormant until invoked by remotely to perform a function
- **Spyware**: inserted into the computer to capture users' data
- **Ransomware**: locks up data via encryption, demanding payment to unlock it

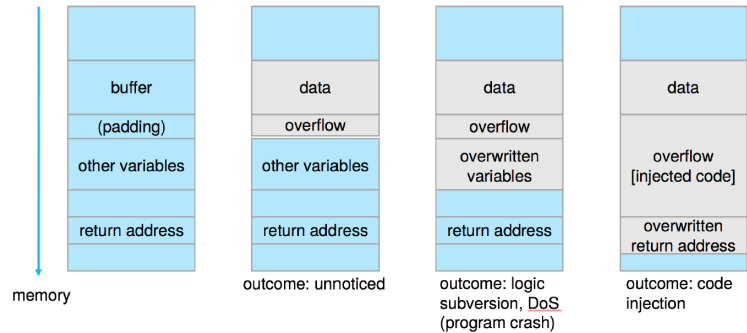
# How to Inject Malware?

## ❑ Code-injection attack

- The system code is not malicious but has bugs allowing executable code to be added or modified


## ❑ Buffer overflow

```
#include <stdio.h>
#define BUFFER SIZE 256
int main(int argc, char *argv[])
{
    char buffer[BUFFER SIZE];
    if (argc < 2)
        return -1;
    else {
        strcpy(buffer, argv[1]);
        return 0;
    }
}
```



# How to Propagate Malware?


- ❑ Virus tries to copy itself to other programs and propagate it to other systems



# How to Hide Malware?

## ❑ Rootkits hook the function called by the anti-malware, and remove their existences

- Windows:
  - NtOpenProcess
  - NtQuerySystemInformation
  - PsActiveProcessLinkHead
- Linux:
  - Proc filesystem (procfs)
  - task\_struct



# Outline

## ❑ System security basis

- Definition
- Security properties

## ❑ Security Protection Stages employed by OS

- Authentication
- Access Control
- Audit

## ❑ Security Threats in OS

## ❑ Hardware Protection


- Integrity verification
- Trusted Execution Environment



# Computer System: A Hierarchic View

## □ System layers


- Different scenarios may have different layers
- Lower layers have higher privileges and can protect higher layers.
- Lower layers need to be better protected



# Chains of Trust

## ❑ Establish verified systems from bottom to top

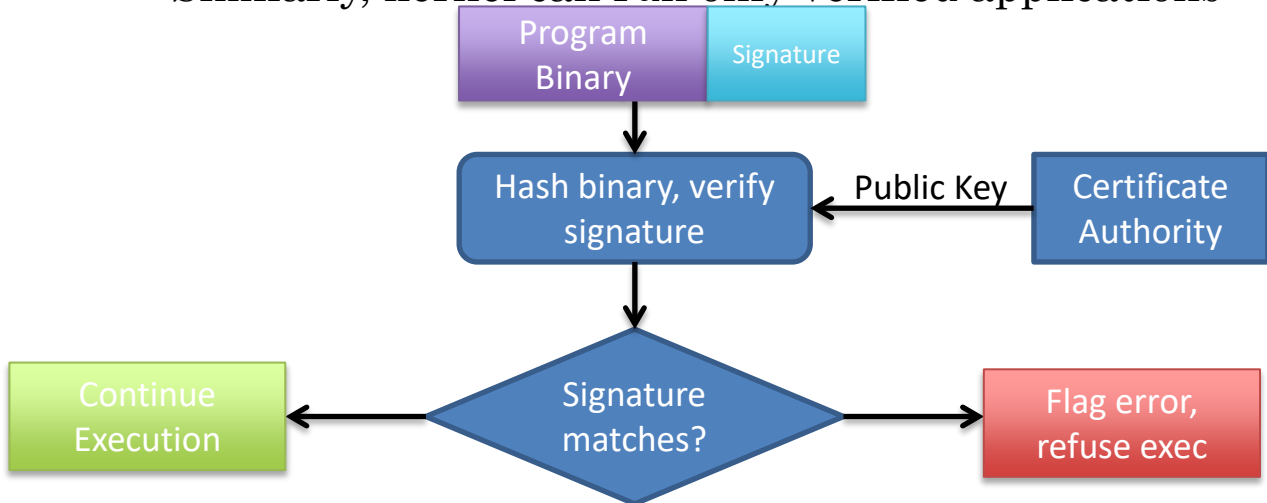
- The bottom layers validate the integrity of the top layers
- If the verification passes, then it is safe to launch it.
- Each layer is vulnerable to attack from below if the lower layers are not secured appropriately



# Integrity Verification

## ❑ Only execute code signed by an entity we trust

- Load the bootloader in the firmware
- Reads and verifies the kernel
- Only loads kernel if the signature is verified
- Similarly, kernel can run only verified applications




# Protect Applications from Untrusted OS

## ❑ Building a secure OS is difficult

- Large code base size and complex functionalities
- An untrusted OS can compromise all applications
- Can we protect the security of apps even when the OS is malicious?

## ❑ Solution: Intel Secure Guard Extension (SGX)

- Security critical code isolated in enclave
- Only CPU is trusted
- Memory is encrypted
- Support remote attestation



# Thank You!

tianwei.zhang@ntu.edu.sg